# When does compressing images pay off?

Compression pays when the input is long and the answer short, or when many requests share the GPU.

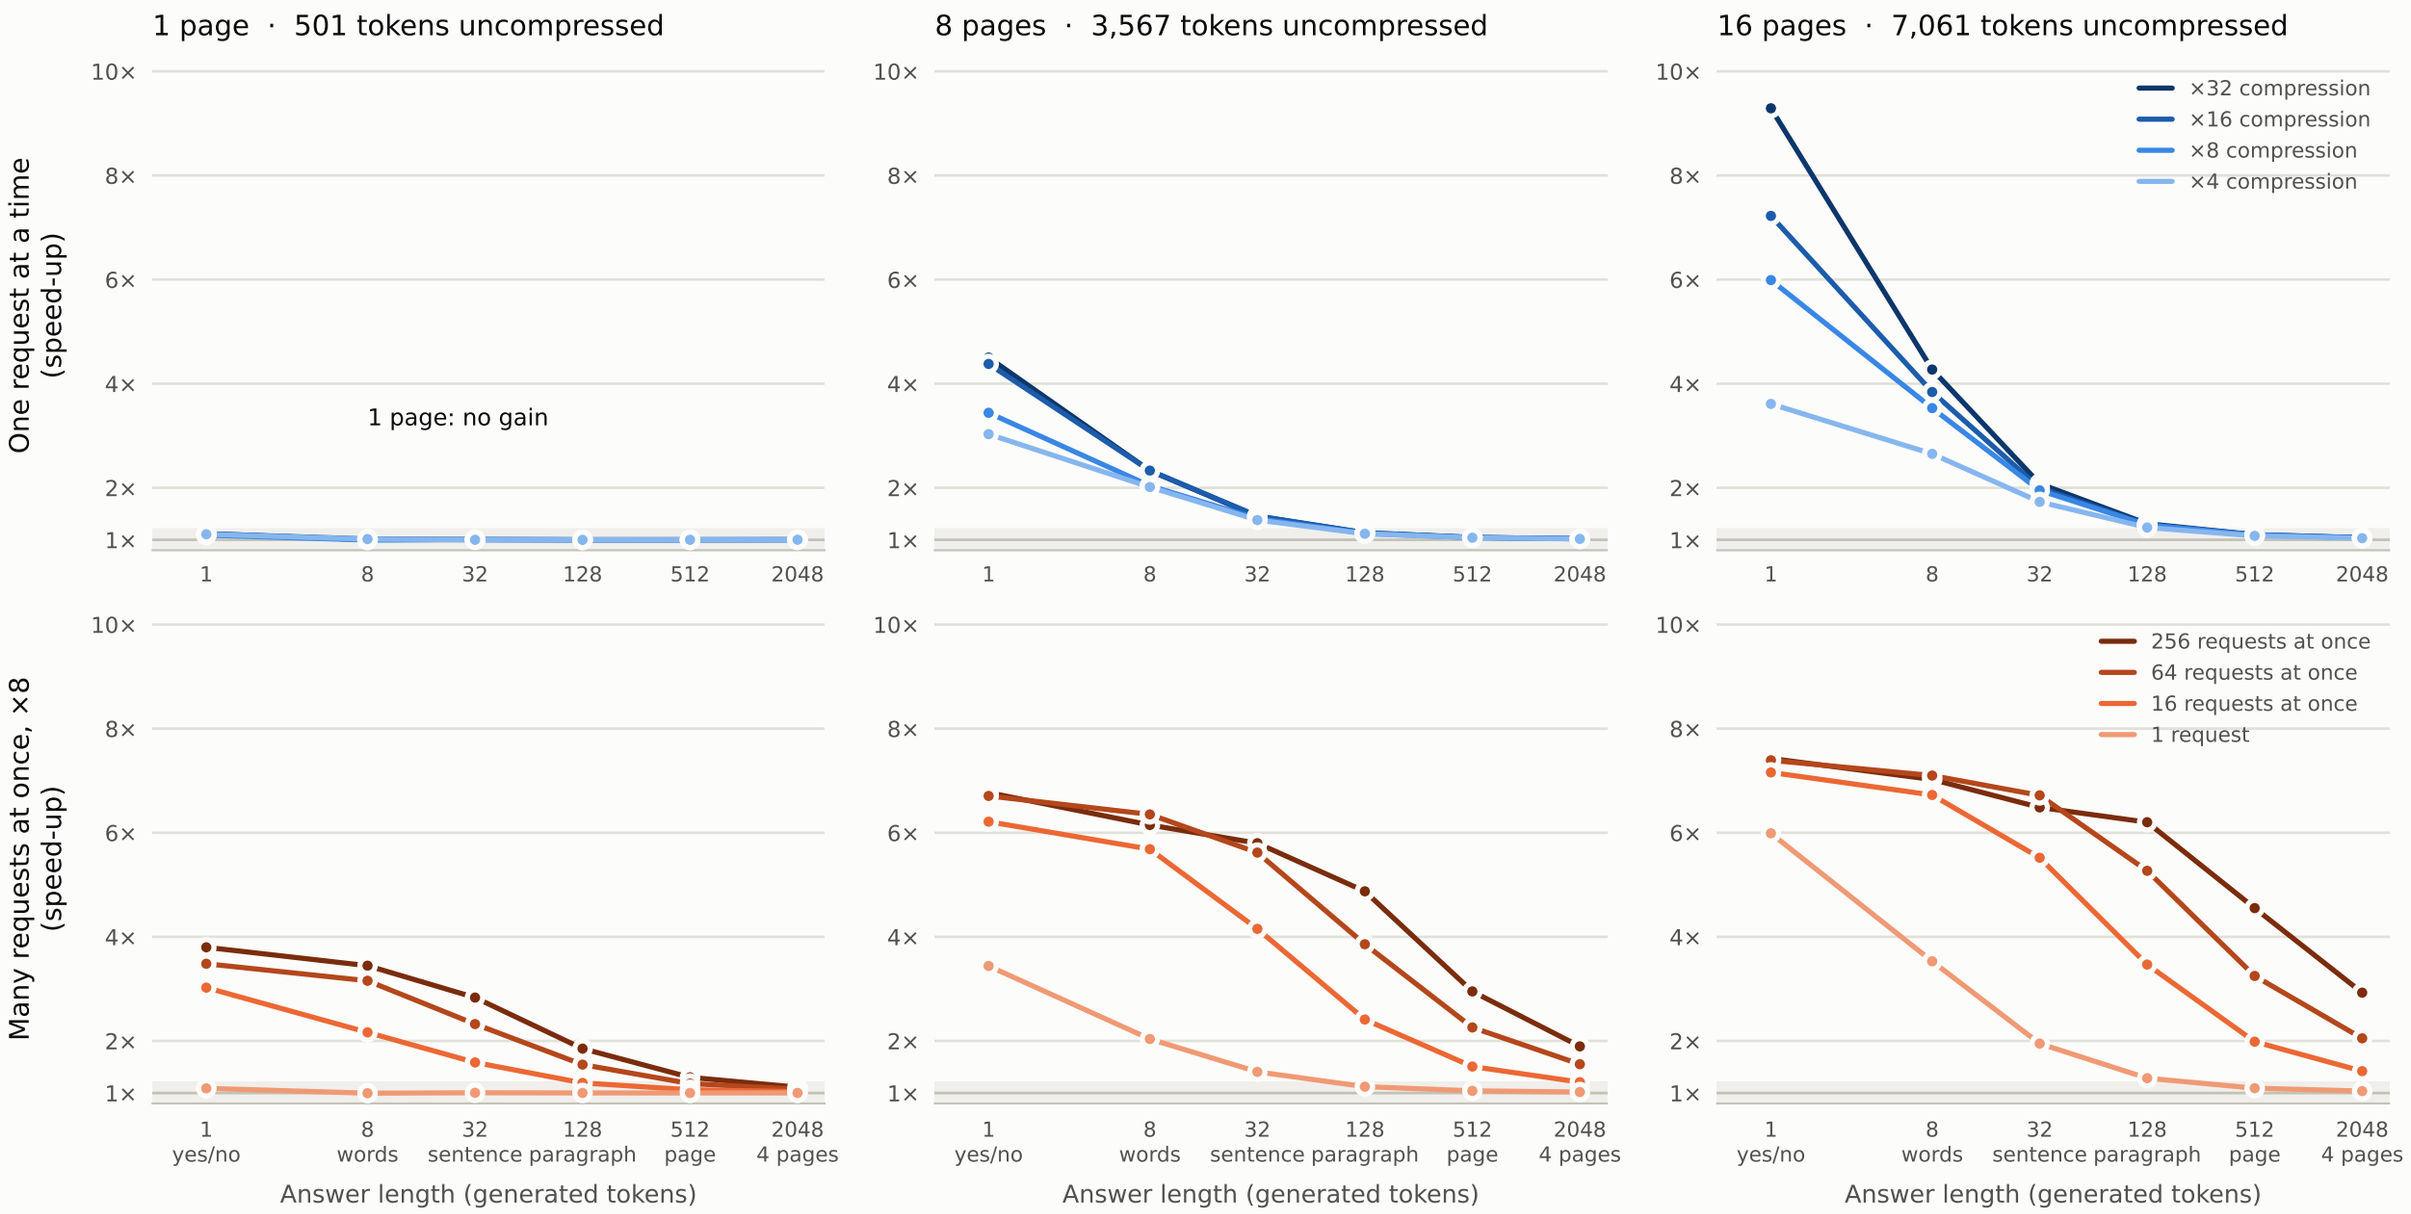


Measured: Qwen3.5-4B in vLLM on one A100; every request has its own prompt, nothing is cached. Grey band: under 1.2× faster, not worth it. Compression ×r keeps 434/r + 1 vectors per page; the text part (63 tokens) is never compressed. Compression runs once, offline (not counted).