When does compressing images pay off?

Compression pays when the answer is short compared to the input.

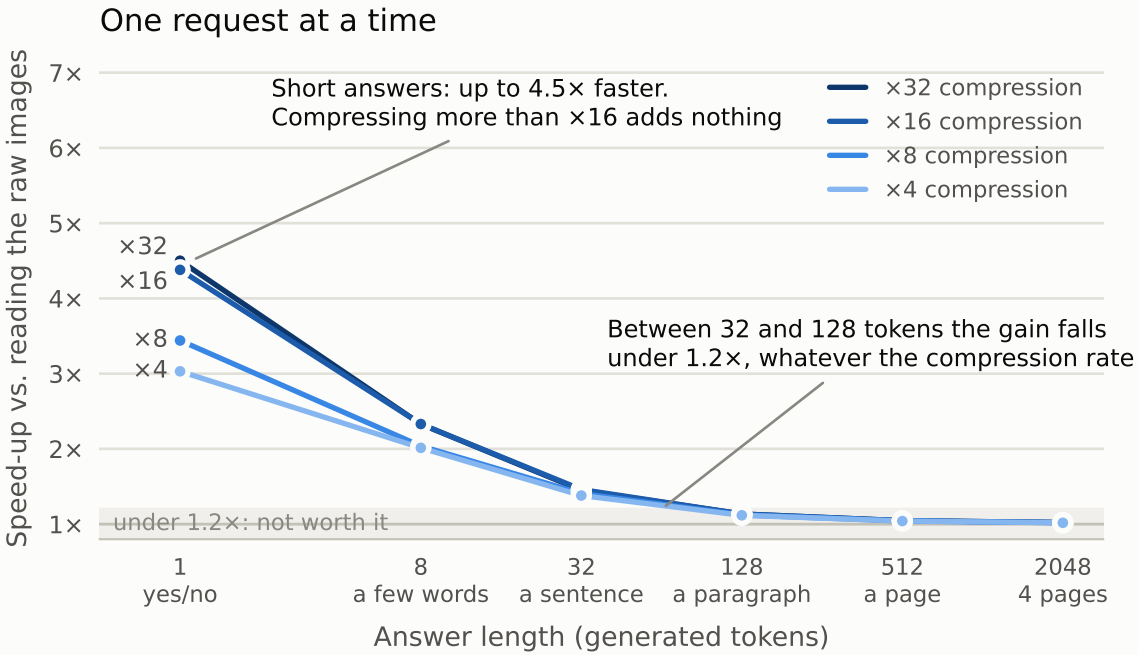

Measured: Qwen3.5-4B in vLLM on one A100; every request has its own prompt, nothing is cached. Input: 8 document pages, 3,567 prompt tokens uncompressed; compression ×r keeps 434/r + 1 vectors per page. Compression runs once, offline (not counted).